

# Why Datacenter KV Cache Systems Fail on HPC

## Datacenter

- ✓ Local NVMe SSD (3-6 GB/s)
- ✓ RDMA / InfiniBand

## HPC (Perlmutter)

- ✗ No Local NVMe
- ✓ Lustre PFS + Slingshot

LMCache/Mooncake

→ 29× slower